Cross-Lingual Anti-Christian Hate Speech

Who Attacks, How, and Why — A Trilingual Computational Analysis

Zhidian Huang

Department of Computational Linguistics · University of Zurich

April 13, 2026

Roadmap

1. Background & Research Questions
2. Data & Augmentation
3. Analytical Pipeline & BERTopic
4. RQ1 — Who? Method & Results
5. RQ2 — How? Method & Results
6. RQ3 — Why? Method & Results
7. Discussion
8. Contributions, Limitations & Outlook

The Problem in One Sentence

"Hate speech research talks a lot about **what** *people say.*

We ask **who** *they attack,* **how**, *and* **why**."

Three blind spots in the literature

- **English-centric** corpora dominate
- **Binary** hate/not-hate is the default task
- **Anti-Christian** discourse in East Asia barely studied

This thesis

Builds a **trilingual** (EN / ZH / JA) interpretable analytical pipeline —

from *detection*
to structured *discourse analysis*

Three Linked Research Questions

Who is targeted?

RQ1

Who?
Target identification

→

What rhetoric is used?

RQ2

How?
Attack frame taxonomy

→

What drives the hostility?

RQ3

Why?
Moral motivations

Case study: **anti-Christian discourse** in Chinese, Japanese & English —
the *same target* appears across very different cultural contexts.

A Trilingual Corpus — Built from Scratch

Corpus at a glance

Lang	Texts	Source
EN	58,940	4 public datasets
ZH	1,965	Tieba (crawled)
JA	2,360	5ch (crawled)
Total	**63,265**	

EN: filtered from HateXplain, Jigsaw, MLMA, Measuring Hate Speech

ZH/JA: manual annotation (Label Studio) + LLM localization (back-translation)

The low-resource challenge

ZH and JA hate speech data is scarce.
Naively-trained classifiers perform very poorly.

Two-stage pipeline for ZH / JA

Stage 1 — Native:

- Crawl → Manual annotation → Fine-tune classifier (Focal Loss)

Stage 2 — Augmentation:

- Translate EN seed texts into ZH/JA
- Filter with the same classifier

Key failure: LLM seed-guided generation collapses into keyword loops after 20–30 iterations
(JA: 100% redundant at 30 samples; ZH: 90% at 50)
⇒ back-translation is more robust.

Data Augmentation: Three Strategies, One Clear Winner

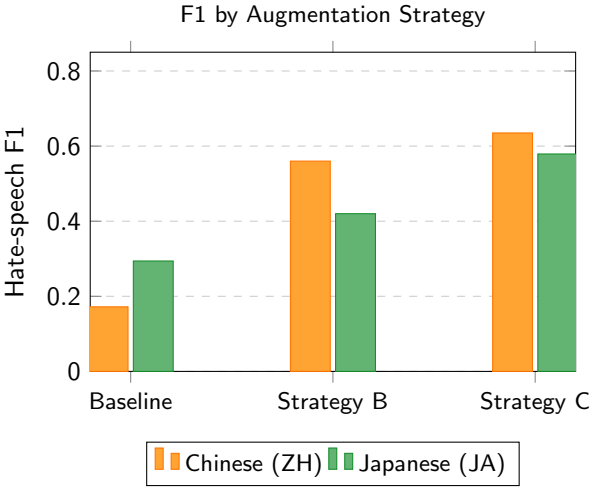

Strategy A — Baseline

Native data only.
Very low F1 due to data scarcity.

Strategy B — Train-only BT

Back-translate training set only.
Distribution shift between train & eval.

✔ Strategy C — Double BT *(Winner)*

Back-translate **both** train & eval sets.
Eliminates distribution gap.

ZH: 0.172 → 0.635
JA: 0.294 → 0.579

The Who–How–Why Pipeline

Input: 63,265 ZH/JA/EN texts

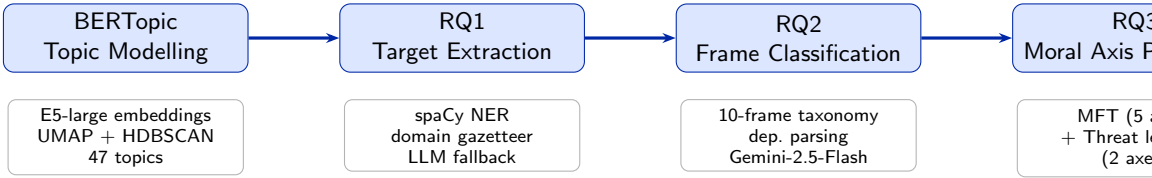

Design principle: thematic structure is *induced* from the corpus (not imposed), so findings are not pre-biased toward English-language research assumptions.

BERTopic: 47 Cross-Lingual Topics

Documents and Topics

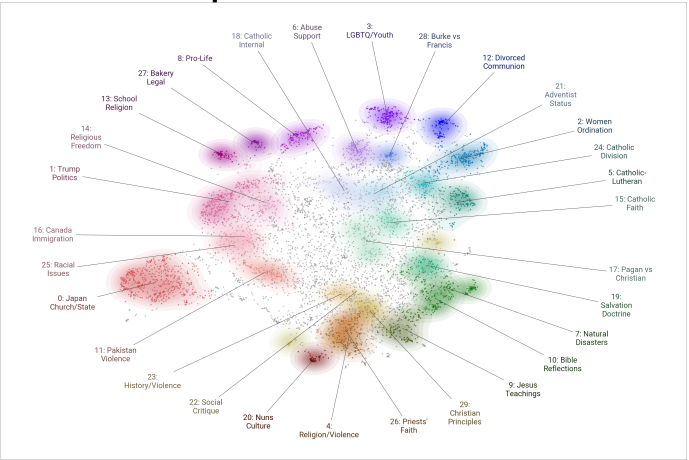

Document map. Each point = one text; colours = topic cluster.

Key numbers

Topics found	**47**
Noise docs	2,530 (4%)
Largest topic	499 docs

Monolingual topics

- JP (99%): Unification Church & Abe
- EN (92%): Religion & Canadian identity
- ZH (96%): Missionaries in China

Cross-lingual (entropy ≈1.1)

- LGBTQ+ & Catholic teaching
- Science vs. creationism
- Conservative/liberal church split

RQ1 — Method: Three-Stage Target Extraction

Topic-labelled corpus (47 topics)

→

Stage 1 — NER

spaCy transformer models
`en_core_web_trf`
`zh_core_web_trf`
`ja_core_news_trf`

Types: PERSON, ORG, NORP, GPE, FAC

→

Stage 2 — Gazetteer

40+ category domain lexicon (clergy roles, institutions, denominational labels)

Catches NER misses

→

Stage 3 — LLM Fallback

Gemini API

For zero-subject constructions in ZH/JA

Async + rate-limit + disk checkpoint

→

if no target found ↓

A **blacklist** removes pronouns and generic religious nouns ("faith", "religion") that appear universally and carry no discriminative information about specific targets.

RQ1 — Results: Top-20 Targets per Language

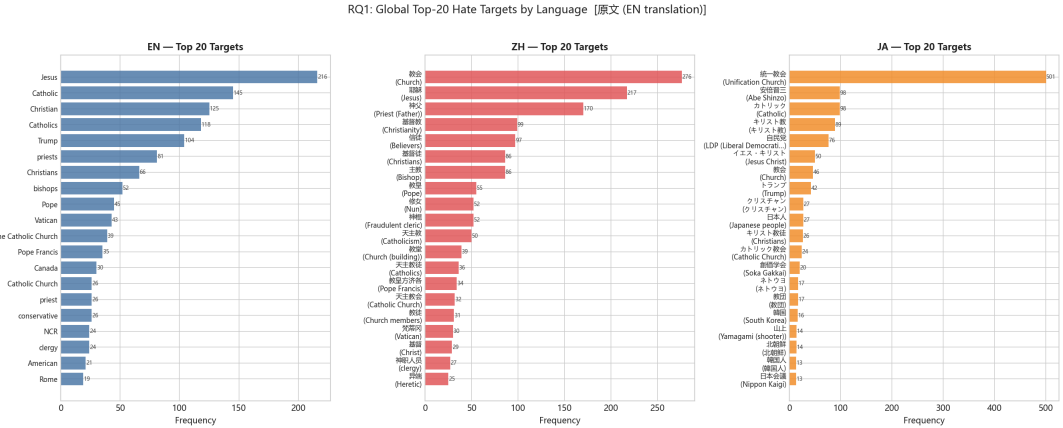

Top-20 hate-speech targets per language, ranked by total mention frequency. Labels: original language + English translation.

RQ1 — Results: Target Profiles Differ by Language

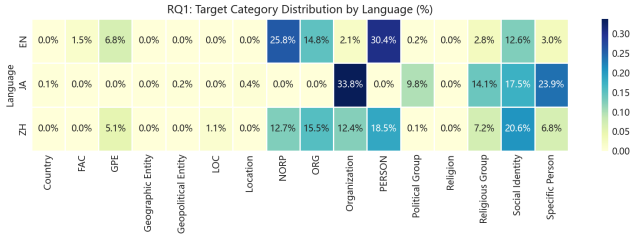

Target-category frequency heatmap by language.

Language-specific patterns

English — *clergy* dominate; named individuals prominent

Chinese — *church institutions* & historical missionaries; state-framed discourse

Japanese — *Unification Church* as focal point; political entanglement salient

Key finding

Attack targets reflect each community's **socio-political context**,
not just shared religious subject matter.

RQ2 — Method: Three-Layer Predicate Extraction

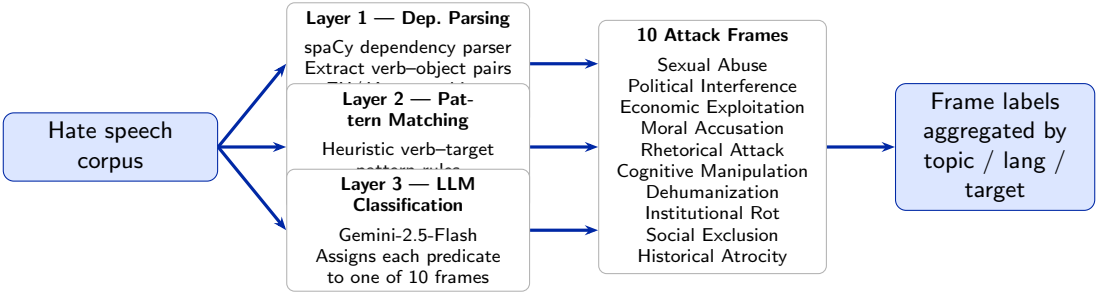

RQ2 — Results: Frame Distribution by Language

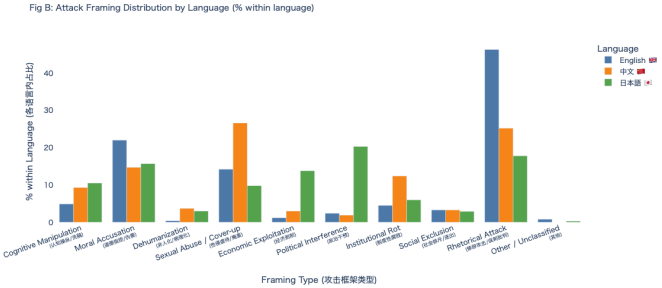

10-frame distribution by language community (predicate counts).

EN — Rhetoric & moral condemnation

Rhetorical Attack (46%),
Moral Accusation (22%),
Sexual Abuse (14%)

JA — Political & economic frames

Political Interference (20%),
Rhetorical Attack (18%),
Economic Exploitation (14%)

Driven by Abe/Unification Church

ZH — Sexual abuse salient

Sexual Abuse Frame (26%),
Rhetorical Attack (25%),
Moral Accusation (14%)

RQ2 — Key Finding: Frame = Target Type, Not Language

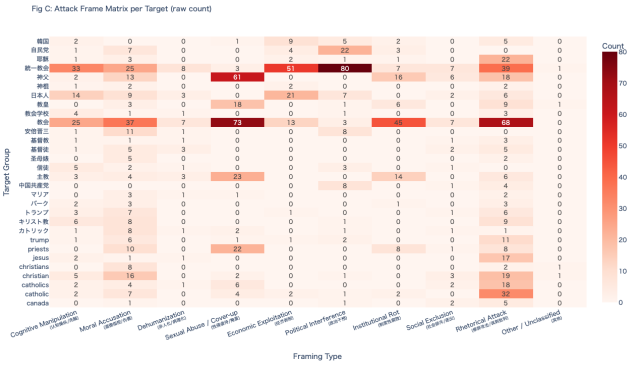

Attack-frame distribution by target entity (all languages pooled).

▲ Structural regularity #1

Frame choice is governed by the **target's social category**, not by language community.

Cross-lingual genre conventions

Target type	Dominant frame
Clergy roles	Sexual Abuse
Church institution	Institutional Rot
Political institution	Political Interf.
Collective identity	Rhetorical Attack
Collective identity	Moral Accusation

Holds independently in EN, ZH & JA.

RQ3 — Method: Moral Axis Projection

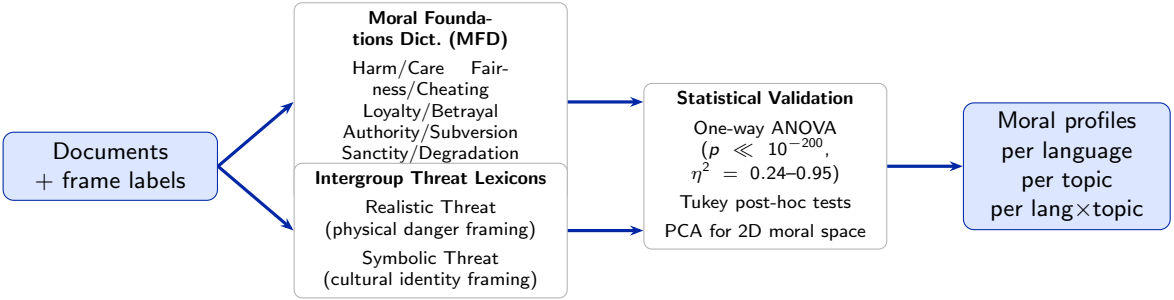

Interpretation: a *negative* score means the discourse invokes the attack-oriented pole of that foundation (e.g., *Harm* rather than *Care*).

RQ3 — Results: Different Moral Vocabularies

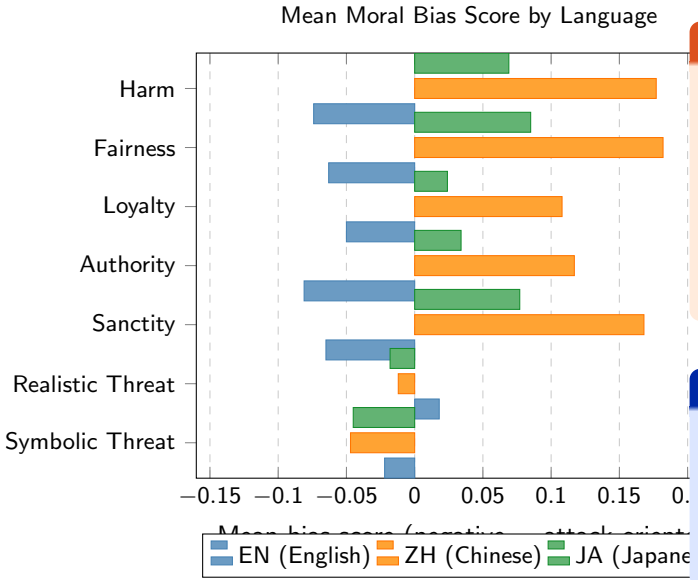

▲ Structural regularity #2

English: negative across all 5 classical MFT axes
⇒ *broad ethical condemnation*

ZH & JA: positive on MFT, negative only on Symbolic + Realistic *Threat*
⇒ *identity-threat framing*

Statistical validation

All 7 axes: ANOVA
$p < 4.9 \times 10^{-324}$
$\eta^2 = 0.24$–0.95

EN is the **only** language where Realistic Threat is *positive* (+0.018).

RQ3 — Results: Two Moral Universes

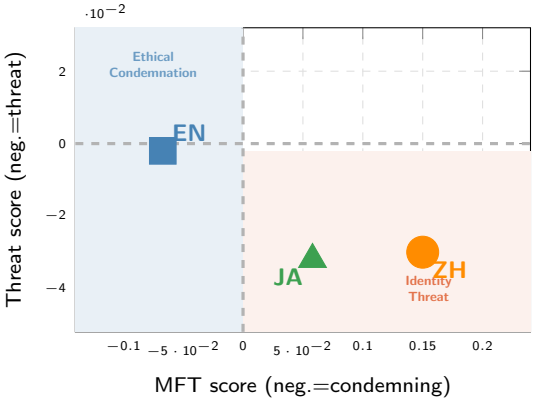

X = mean of 5 MFT axes; Y = mean of 2 threat axes. All differences significant ($p < 10^{-200}$).

English speakers say:

"Christianity is *morally wrong*"

Wide ethical condemnation using *shared* normative standards (insider critique)

East Asian speakers say:

"Christianity is *dangerous to us*"

Cultural invasion / social threat framing (outsider critique, group safety)

Same target. Fundamentally different moral logic.

Discussion: Two Structural Regularities

#1 — Attack Frames as Genre Conventions

Frame choice is governed by **target type**, not language.
Same conventions hold in EN, ZH & JA.

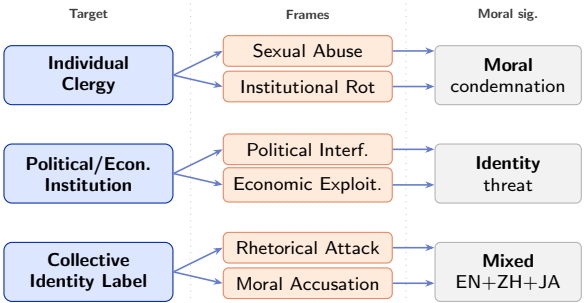

Pattern stable across EN, ZH & JA

#2 — Two Moral Architectures

EN: *moral transgressor* model
— Christianity judged by universal ethical standards.

ZH/JA: *outgroup threat* model
— Christianity is a cultural/social danger.

Binary classifiers collapse this difference; structured analysis makes it visible.

Broader lesson

Binary detection cannot surface these structural patterns.
The who–how–why pipeline can.

Discussion: Methodological Reflections

What worked

- **Double BT** resolves distribution shift — reusable solution for low-resource hate speech
- **Cross-lingual centroid alignment** in BERTopic creates genuinely mixed topics without translation
- **Corpus-induction** lets culturally specific topics emerge naturally (JP Unification Church)

What to watch for

- **LLM data collapse**: after 20–30 iterations, ZH/JA generation becomes 90–100% redundant
- **Platform bias**: Tieba (ZH) and 5ch (JA) are not representative of all online speech
- **Annotation noise**: implicit hate labels have moderate inter-annotator agreement

Contributions

Resource

First trilingual corpus of anti-Christian hate speech (EN/ZH/JA, 63,265 texts)

Structured annotation: target $\times$ frame $\times$ moral axis

Open-source on GitHub

Empirical

First systematic cross-lingual analysis revealing:

(1) Attack frame = genre convention

(2) EN vs. EA communities inhabit different moral universes

Methodological

Reusable **who–how–why** pipeline:

multilingual embeddings,
topic modelling,
frame classification,
moral-axis projection

All code open-source

Limitations & Future Work

Current limitations

- **Platform coverage**: Tieba & 5ch may not generalise to WeChat, Twitter/X, LINE
- **Temporal snapshot**: discourse evolves — longitudinal re-running needed
- **ZH/JA corpus size**: still small despite augmentation
- **Annotation noise**: implicit hate labels are harder to label consistently

Promising directions

- Apply pipeline to **other religious minorities** (Islamophobia, antisemitism) — same framework, different case
- **More languages** (Korean, Vietnamese): does the moral-architecture split generalise?
- **Larger LLMs with RAG** for stable low-resource augmentation
- **Longitudinal analysis**: does moral framing shift after major events? (post-2022 Abe assassination)

Summary: The Story in Three Acts

Act 1 — Data

Built a trilingual corpus from near-zero.

Back-translation solved the low-resource problem better than LLM generation.

Act 2 — Structure

Frame selection follows the target, not the language.

Cross-cultural genre conventions exist in hate speech.

Act 3 — Morality

English condemns.
East Asian fears.

Same target, fundamentally different moral logic — invisible to binary detectors.

`github.com/notsofun/Master_Thesis`

Open-source corpus, annotations, and pipeline code

Thank you.

Questions?

Zhidian Huang
huangzhidian15@gmail.com

github.com/notsofun/Master_Thesis

Selected References I